

RecruitGPT X

Ranking candidates like a great recruiter — not a keyword filter

India Runs Data & AI Challenge · Redrob Senior AI Engineer

The Problem

- Recruiters review hundreds of profiles but keyword filters miss real fit.
- Honeypot profiles stuff AI skills into unrelated titles (HR Manager, Accountant).
- Behavioral signals — response rate, GitHub, recruiter saves — are ignored by ATS.

What the Role Needs

- Parsed job_description.docx into jd_config.py: embeddings, retrieval, ranking, vector DBs.
- 5–9 years experience, India location, production deployment language.
- Redrob behavioral modifier from redrob_signals_doc.

Hybrid Offline Ranker

- Weighted: title 32%, skills 28%, production 18%, experience 12%, location 10%.
- Modifiers: consulting-only penalty, behavioral boost (capped), honeypot trap.
- CPU-only, no network — 100K candidates in ~7 seconds.

Honeypot Defense

- Weak titles capped; noise skills (HTML, Tailwind) penalized.
- Unrelated title + inflated AI skills → 0.25× demotion.
- Zero honeypot titles in our top-100 shortlist.

Results

- #1 CAND_0002025 — Senior AI Engineer, 7 core IR skills, 0.80 response rate.

- #2 CAND_0071974 — Senior AI Engineer, 8 core IR skills.
- Reproduce: `python rank.py --candidates ./data/candidates.jsonl --out ./submission.csv`

Live Platform (RecruitGPT X)

- Next.js + FastAPI + LangGraph 7-agent pipeline for interactive recruiting.
- Semantic matching, explainability, bias audit, Candidate Radar, skill evolution.
- Same philosophy as offline ranker: understand fit, explain every rank.